


Semi-supervised learning (SSL)

for remote sensing (SSL-for-RS)



About me




- Itzá Hernández
- PhD. Computer Science - Computer Vision
- University Jaume I, Castellón, Spain
- Internship Helmholtz Institute Freiberg for Resource Technology (HIF)
- @itzahs   

Post-it



- **Name:** Itzá
- **First Job:** Institute for Tropical Agriculture
- **What I learned from it:** If you plant coffee and cacao together, your coffee beans will have a pinch of cacao flavor.

Learning outcomes of this workshop



Content of this workshop



01

Introduction to DSSL

16:00 - 16:15 (15 min)

02

Data & Code

16:15 - 16:30 (15 min)

03

Training model

16:30 - 17:00 (30 min)

04

Evaluate model

17:00 - 17:30 (30 min)

Questions welcome anytime

Total estimated time:
≈ 2 hours



01 Introduction

Deep Semi Supervised Learning
16:00 - 16:15 (15 min)




Machine Learning & Deep Learning

ML



Feature extraction




Classification



Class probability

DL



Feature extraction + Classification



Class probability


Supervised

Semi-supervised

Unsupervised

DL in Aerial Image Classification


Pixel-level



Object-level





Scene-level



Convolutional Neural Network (CNN)



Original Image



Taxonomy


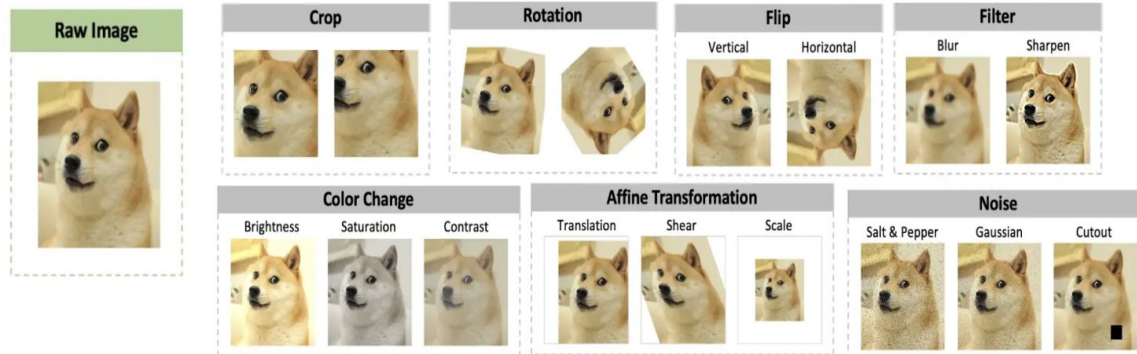


Image data augmentation for DL



- **Data Warping:** transforms existing images such as the label is preserved.

neutral



Generated images

fear



angry




Synthetic data - CycleGANS - emotion classification

- **Oversampling:** creates synthetic instances and add them to the training set.

[4] Sharma, A. (2019, June 12). Complete Guide to Data Augmentation for Computer Vision. Towards Data Science. <https://towardsdatascience.com/complete-guide-to-data-augmentation-for-computer-vision-1abe4063ad07>.

[5] Shorten, C., Khoshgoftaar, T.M. A survey on Image Data Augmentation for Deep Learning. *J Big Data* 6, 60 (2019). <https://doi.org/10.1186/s40537-019-0197-0>


Generative methods



Generative modeling refers to the practice of **creating artificial instances** from a dataset such that they retain similar characteristics to the original dataset.

e.g. Generative Adversarial Networks (**GAN**) and Variational AutoEncoder (**VAE**).



Consistency regularization




A consistency regularization term is applied to the **final loss function**. e.g. **Cross Entropy for example $H(F(x), T_x)$** .

A realistic perturbation in the training data should not change the output of the model.

e.g. Unsupervised Data Augmentation (**UDA**).

 = Basic Neural Network Layer  = Augmentation Operator.
RandAugment and Back-translation for UDA.

Graph-based




The nodes/vertices are **representations of the training samples** and the edges encode the relationships between the nodes.

The goal is to encode the nodes as small-scale vectors at first and then how each node belongs within the context in the graph.

Deep embedding methods are **AutoEncoders** and Graph Neural Networks (GNN).

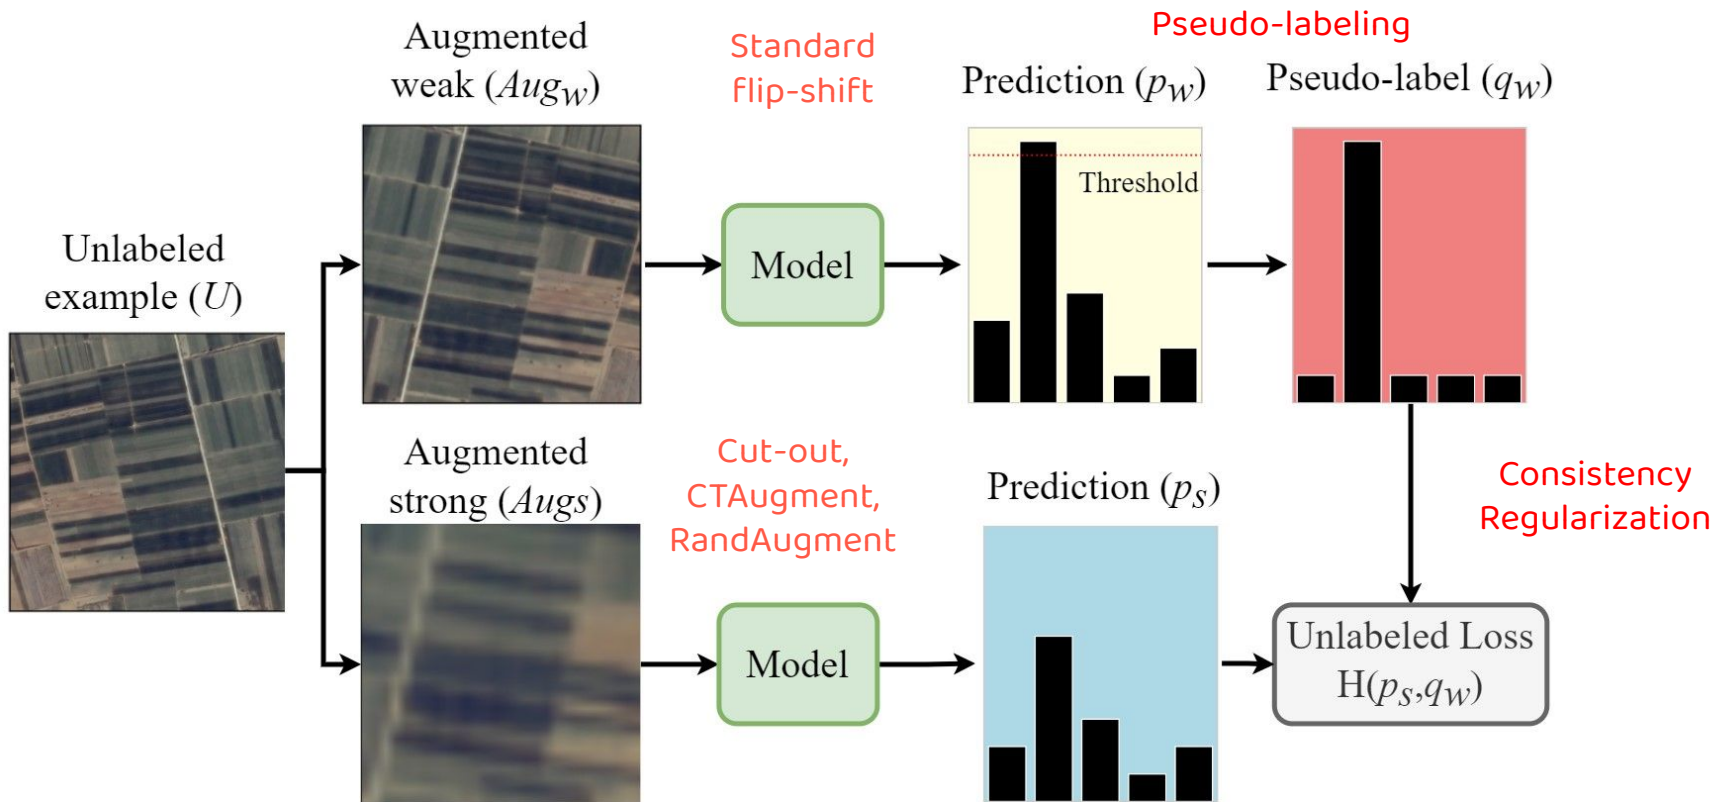
Pseudo-labelling



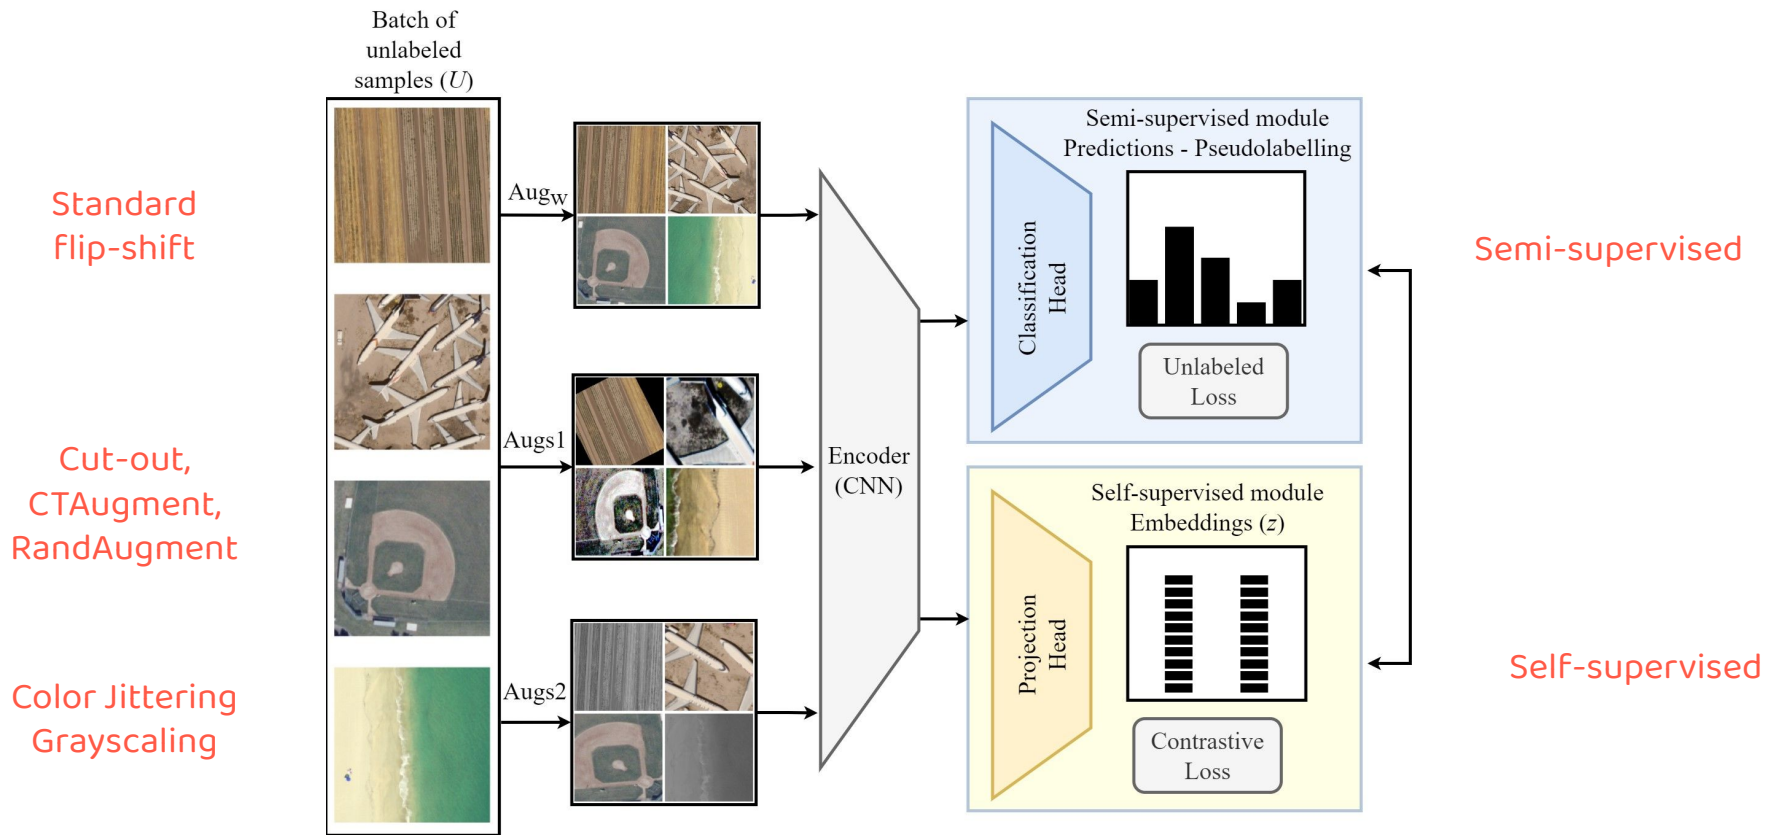
Pseudo-labeling methods rely on the **high-confidence of pseudo-labels**, which can be added to the training data set as labeled data.

Self-training: leverage model's own confident predictions to produce the pseudo-labels for **unlabeled data**.

Hybrid: Fixmatch - Consistency Regularization and Pseudo Labeling



Hybrid: CoMatch - Consistency Regularization and Pseudo Labeling



02

Data & Code

Cloning the GitHub repository and
downloading the dataset
16:15 - 16:30 (15 min)



1 or 2



- Do you work 1- with satellite imagery or 2- other types of data.
- When using Deep Learning do you prefer working with 1-Pytorch or 2-Tensorflow.
- For working you use 1-Cloud services or 2-Local servers/Personal laptop.



Codes for this workshop



@itzahs **SSL-for-RS**

SSL-for-RS

Public

Configuration files for the ICIAP2023 paper "Semi-supervised classification for remote sensing datasets" & step-by-step implementation for the Geomundus2023 Workshop.

 Jupyter Notebook



Step-by-step implementation
in 4 notebooks on Colab 



Visual Studio Code

SCAN ME

Semi-supervised learning toolbox

Class-Aware Contrastive Semi-Supervised Learning

Publisher: IEEE

[Cite This](#)

[PDF](#)

Fan Yang ; Kai Wu ; Shuyi Zhang ; Guannan Jiang ; Yong Liu ; Feng Zheng ; Wei Zhang ; Chengjie Wang ; Long Zeng [All Authors](#)

5

Paper
Citations

142

Full
Text Views



Abstract

Document Sections

1. Introduction
2. Related Work
3. Method
4. Experiments
5. Ablation

[Show Full Outline](#)

[Authors](#)

[Figures](#)

[References](#)

[Citations](#)

Abstract:

Pseudo-label-based semi-supervised learning (SSL) has achieved great success on raw data utilization. However, its training procedure suffers from confirmation bias due to the noise contained in self-generated artificial labels. Moreover, the model's judgment becomes noisier in real-world applications with extensive out-of-distribution data. To address this issue, we propose a general method named Class-aware Contrastive Semi-Supervised Learning (CCSSL), which is a drop-in helper to improve the pseudo-label quality and enhance the model's robustness in the real-world setting. Rather than treating real-world data as a union set, our method separately handles reliable in-distribution data with class-wise clustering for blending into downstream tasks and noisy out-of-distribution data with image-wise contrastive for better generalization. Furthermore, by applying target reweighting, we successfully emphasize clean label learning and simultaneously reduce noisy label learning. Despite its simplicity, our proposed CCSSL has significant performance improvements over the state-of-the-art SSL methods on the standard datasets CIFAR100 [18] and STL10 [8]. On the real-world dataset Semi-iNat 2021 [27], we improve FixMatch [25] by 9.80% and CoMatch [19] by 3.18%. Code is available <https://github.com/TencentYoutuResearch/Classification-SemiCLS>.

Published in: 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)

Date of Conference: 18-24 June 2022

INSPEC Accession Number: 22095223

Date Added to IEEE Xplore: 27 September 2022

DOI: 10.1109/CVPR52688.2022.01402

Semi-supervised learning toolbox



Tencent YouTu Research

119 followers Shanghai, China

Classification-SemiCLS

Public

Code for CVPR 2022 paper "Class-Aware Contrastive Semi-Supervised Learning"

Python 60 11 4 0 Updated on Dec 19, 2022


- configs
- dataset
- loss
- models
- optimizer
- pretrained_models
- scheduler
- tools
- trainer
- utils

<https://github.com/TencentYoutuResearch/Classification-SemiCLS>

Dataset

UCM [7] Dataset [[download](#)]:


- 21 clases
- 100 images per class
- 256 x 256 pixels (0.3 m)
- Released 2010



agricultural




airplane



baseball diamond



dense residential



forest



free-way



medium residential



mobile home park



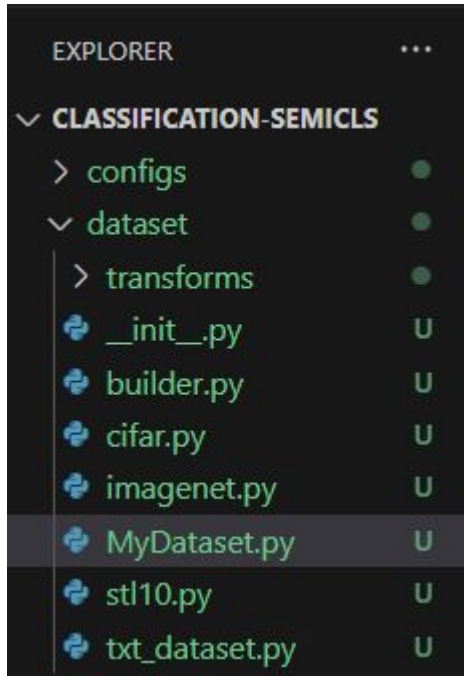
overpass

Sampling strategy

- 50 % training and 50% testing (Cheng, G. et al, 2020) [2].
- When using supervised 1,050 samples for training and testing.
- When using semi-supervised: 4 labels per class or 84 labeled samples in total for training (Sohn, K. et al, 2020) [6].
- Labeled to unlabeled ratio: 1-7
 - Depending on the batch size, if Batch labeled=8 then Batch unlabeled=56.

Dataset	Class	# Labels per class	# Training data	# Testing data	Train/Test split
UCM	21	4/25/40	84/525/840	1,050	50% / 50%

Creating train and test .txt



```
#val, 1-train dataset
class MyDataset(Dataset):
    """
    Interface provided for customized data sets

    names_file: a txt file, each line in the form of "image_path label"

    transform: transform pipeline for mydataset
    """
```

/content/MyDrive/SSL4RS/Classification-SemiCLS/data/UCM/Images

UCM_test.txt

UCM_train.txt

```
./data/UCM/Images/agricultural/agricultural70.tif 0
./data/UCM/Images/agricultural/agricultural59.tif 0
```


Transformation pipeline - No Augmentation

```
names_file = './data/UCM/Images/UCM_train.txt'  
dataset = MyDataset(names_file, transform=None)
```

0



3



Transformation pipeline - Weak Augmentation

```
transform_w = transforms.Compose([
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.Resize(size=256),
    transforms.CenterCrop(size=224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])])
dataset_w = MyDataset(names_file, transform=transform_w)
```

torchvision.transforms




Transformation pipeline - Strong Augmentation 1

```
from dataset.transforms.randaugment import RandAugmentMC
transform_s1 = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomResizedCrop(size=224, scale=(0.2, 1.0)),
    transforms.Lambda(lambda img: RandAugmentMC(n=2, m=10)(img)), # Use RandAugmentMC directly with desired parameters
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225))
])
dataset_s1 = MyDataset(names_file, transform=transform_s1)
```

`Dataset.transforms.randaugment`

`Class RandAugmentMC`




Transformation pipeline - Strong Augmentation 2

```
transform_s2 = transforms.Compose([
    transforms.RandomResizedCrop(size=224, scale=(0.2, 1.0)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomApply([
        transforms.ColorJitter(brightness=0.4, contrast=0.4, saturation=0.4, hue=0.1)], p=0.8),
    transforms.RandomGrayscale(p=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406), std=(0.229, 0.224, 0.225))
])
dataset_s2 = MyDataset(names_file, transform=transform_s2)
```

0



3

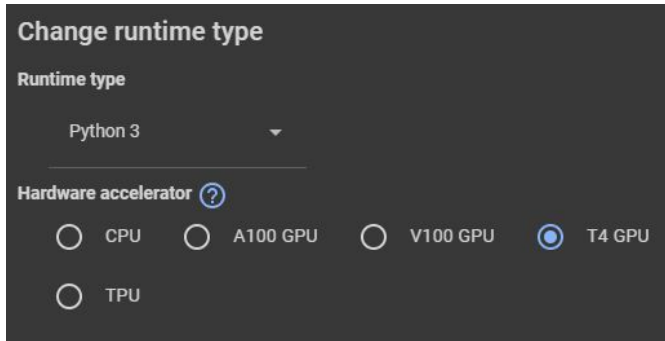


Model training

Set up environment, adapt the configuration
files and launch the training
16:30 - 17:00 (30 min)



Checking GPU availability



GPU's memory usage in MiB
(mebibytes) 0MiB / 15360MiB.

≈ 16 GB RAM

NVIDIA-SMI 525.105.17 Driver Version: 525.105.17 CUDA Version: 12.0									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr. ECC		
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.	MIG M.	
0	Tesla T4		Off	00000000:00:04.0	Off			0	
N/A	40C	P8	9W / 70W	0MiB / 15360MiB		0%	Default	N/A	
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
No running processes found									

Modifications: dataset/builder.py

```
81     else:
82
83         # check if .ipynb_checkpoints is in root, and exclude it
84         root = os.path.join(cfg.root, '')
85         if os.path.isdir(os.path.join(root, ".ipynb_checkpoints")):
86             exclude_dir = os.path.join(root, ".ipynb_checkpoints")
87         else:
88             exclude_dir = None
```

Google colab creates temporary .ipynb inside the image folders and the dataset class gives error when other than .tif extensions are found.

Modifications: train_semi.py

```
444 data_x = labeled_iter.next()
445 #data_x = next(labeled_iter)
446 except Exception:
447     if args.world_size > 1:
448         labeled_epoch += 1
449         labeled_trainloader.sampler.set_epoch(labeled_epoch)
450         labeled_iter = iter(labeled_trainloader)
451         data_x = labeled_iter.next()
452     #data_x = next(labeled_iter)
453
454 try:
455     data_u = unlabeled_iter.next()
456     #data_u = next(unlabeled_iter)
457 except Exception:
458     if args.world_size > 1:
459         unlabeled_epoch += 1
460         unlabeled_trainloader.sampler.set_epoch(unlabeled_epoch)
461         unlabeled_iter = iter(unlabeled_trainloader)
462         data_u = unlabeled_iter.next()
463     #data_u = next(unlabeled_iter)
```

```
from mmcv import Config
```

```
from mengine.config
import Config
```

Modifications: configuration file

```
train = dict()
```



Specifies algorithm (FixMatch or CoMatch), epochs and loss function

```
model= dict()
```

Details about model architecture (ResNet, WideResNet)

```
data = dict()
```

Parameters for loading and preprocessing the dataset, including number of labeled samples, batch size and augmentation pipeline.



Modifications: configuration file

```
train = dict(eval_step=1024,  
             total_steps=2**20, #1024*20,  
             trainer=dict(type="FixMatch",  
                           threshold=0.95,  
                           T=1.,  
                           lambda_u=1.,  
                           loss_x=dict(  
                               type="cross_entropy",  
                               reduction="mean"),  
                           loss_u=dict(  
                               type="cross_entropy",  
                               reduction="none"),  
                           ))  
num_classes = 10 #21
```

1. How often model performance is evaluated.
2. Overall duration of the training process.

Creating:

```
./Classification-SemiCLS/  
configs/fm_ucm.py
```

Modifications: configuration file

```
model = dict(  
    type="wideresnet",  
    depth=28,  
    widen_factor=2,  
    dropout=0,  
    num_classes=num_classes,  
)
```

```
cifar10_mean = (0.4914, 0.4822, 0.4465)  
cifar10_std = (0.2471, 0.2435, 0.2616)  
  
#ucm_mean = (0.485, 0.456, 0.406)  
#ucm_std = (0.229, 0.224, 0.225)
```

WideResNet - Number of filters in the ResNet increased by a value of 2.

The wider the network, the more filters, the more GPU.

ImageNet mean and std of the pixel intensities for each color channel (RGB).

Modifications: configuration file

```
data = dict(  
    # CIFAR10SSL, CIFAR100SSL  
    type="CIFAR10SSL", # "MyDataset",  
    num_workers=4, #0,  
    num_labeled=250, #84,  
    num_classes=num_classes,  
    batch_size=64, #8,  
    expand_labels=False,  
    mu=7,  
    root="./data/CIFAR",  
    #root="./UCMerced_LandUse/Images",  
    #labeled_names_file="./UCMerced_LandUse/Images/UCM_train.txt",  
    #test_names_file="./UCMerced_LandUse/Images/UCM_test.txt",
```

4 labeled
examples per
class.

Path to images

Modifications: configuration file

```
lpipelines=[  
    # 50% chances that the image is horizontally flipped  
    dict(type="RandomHorizontalFlip"),  
    # RandomCrop crops a fixed size whereas RandomResizedCrop crops and then resizes.  
    dict(type="RandomCrop",  
        size=32,  
        padding=int(32 * 0.125),  
        padding_mode='reflect'),  
    #dict(type="RandomResizedCrop", size=224, scale=(0.2, 1.0)),  
    dict(type="ToTensor"),  
    dict(type="Normalize", mean=cifar10_mean, std=cifar10_std)  
    #dict(type="Normalize", mean=ucm_mean, std=ucm_std)  
],
```

```

upipeline=[
    dict(type="RandomHorizontalFlip"),
    dict(type="RandomCrop",
        size=32,
        padding=int(32 * 0.125),
        padding_mode='reflect'),
    #dict(type="Resize", size=256),
    #dict(type="CenterCrop", size=224),
    dict(type="ToTensor"),
    dict(type="Normalize", mean=cifar10_mean, std=cifar10_std)
],
[
    dict(type="RandomHorizontalFlip"),
    dict(type="RandomCrop",
        size=32,
        padding=int(32 * 0.125),
        padding_mode='reflect'),
    #dict(type="RandomResizedCrop", size=224, scale=(0.2, 1.0)),
    dict(type="RandAugmentMC", n=2, m=10),
    dict(type="ToTensor"),
    dict(type="Normalize", mean=cifar10_mean, std=cifar10_std)
    #dict(type="Normalize", mean=ucm_mean, std=ucm_std)
],

```

Modifications: configuration file

Augmentation strategy

- **Weak:**
 - Flip
 - Crop
- **Strong1:**
 - RandomAugment
 - CutOut
- **Strong2:**
 - Random Color Jittering
 - Grayscale Conversion

Modifications: configuration file

```
vpipeline=[  
    #dict(type="Resize", size=256),  
    dict(type="ToTensor"),  
    dict(type="Normalize", mean=cifar10_mean, std=cifar10_std)  
    #dict(type="Normalize", mean=ucm_mean, std=ucm_std)  
])
```

Some UCM samples have inconsistent spatial sizes, while most have a shape of 256*256*3 some are 253*256*3.

Launch training

```
torch.cuda.OutOfMemoryError: CUDA out of memory. Tried
to allocate 368.00 MiB (GPU 0; 14.75 GiB total
capacity; 14.45 GiB already allocated; 168.81 MiB free;
14.46 GiB reserved in total by PyTorch) If reserved
memory is >> allocated memory try setting
max_split_size_mb to avoid fragmentation. See
documentation for Memory Management and
PYTORCH CUDA ALLOC CONF
```


Not all **trainers** use same amount of GPU

FixMatch
12-24 GB

CoMatch
24-40 GB

+CCSSL
>40 GB

- Image size: 224x224
- Batch size: 8
- WideResNet
- Labeled to unlabeled ratio: 1-7


Launch training


```
2023-10-18 22:14:41,826 - INFO - root - Train Epoch: 1/ 20. Iter: 1/ 5. LR: 0.0300 batch_time: 17.171 data_time: 14.079 loss: 3.103 loss_x: 3.103
2023-10-18 22:14:56,879 - INFO - root - Train Epoch: 1/ 20. Iter: 2/ 5. LR: 0.0300 batch_time: 16.112 data_time: 13.820 loss: 3.100 loss_x: 3.100
2023-10-18 22:15:10,361 - INFO - root - Train Epoch: 1/ 20. Iter: 3/ 5. LR: 0.0300 batch_time: 15.236 data_time: 13.206 loss: 2.877 loss_x: 2.877
2023-10-18 22:15:24,654 - INFO - root - Train Epoch: 1/ 20. Iter: 4/ 5. LR: 0.0300 batch_time: 15.000 data_time: 13.081 loss: 3.006 loss_x: 3.006
2023-10-18 22:15:36,954 - INFO - root - Train Epoch: 1/ 20. Iter: 5/ 5. LR: 0.0299 batch_time: 14.460 data_time: 12.625 loss: 2.953 loss_x: 2.953
Test Iter: 263/ 263. Data: 0.139s. Batch: 0.167s. Loss: 18.7652. top1: 5.24. top5: 23.52. : 100% 263/263 [00:43<00:00, 5.99it/s]
2023-10-18 22:16:20,872 - INFO - root - Epoch 0 top-1 acc: 5.24
2023-10-18 22:16:20,873 - INFO - root - Epoch 0 top-5 acc: 23.52
2023-10-18 22:16:21,043 - INFO - root - Best top-1 acc: 5.24
2023-10-18 22:16:21,043 - INFO - root - Mean top-1 acc: 5.24
```


Open questions:


- What does Test Iter mean?
- Top-1?
- Top-5?
- Best top-1 vs. top-1?

Training output

 2023-10-20_08:19:55_-1.log

 checkpoint.pth.tar

 fm_ucm.py

 model_best.pth.tar

Your directory structure should now resemble the following:

Classification-SemiCLS

results

└ fm_ucm_wres_x2_b1x8_l840

├── fm_ucm_wres_x2_b1x8_l840.log

├── fm_ucm_wres_x2_b1x8_l840.tar

└── fm_ucm_wres_x2_b1x8_l840.py

- The `.py` file represents the configuration file.
- The `.log` file contains the training output.
- The `.tar` file contains the trained model.

- The checkpoint is the last epoch and model_best is the epoch that performed the best on the testing set so far.

Modifications: configuration file

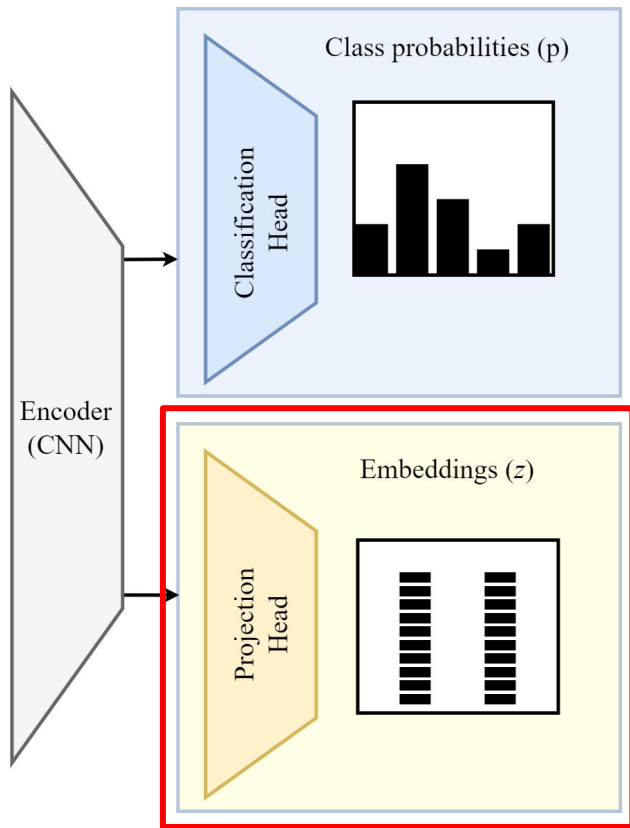
```
train = dict(  
    eval_step=1024,  
    total_steps=2**10*512,  
    trainer=dict(  
        type='CoMatch',  
        threshold=0.95, #pseudolabel threshold  
        queue_batch=5, #memory buffer  
        contrast_threshold=0.8, #similarity matrix  
        da_len=32, #distribution alignment  
        T=0.2, # temperature  
        alpha=0.9, # 1-alpha for memory smoothed pseudo label  
        lambda_u=1.0, #unlabeled loss  
        lambda_c=1.0, #contrastive loss  
        #supervised loss  
        loss_x=dict(type="cross_entropy", reduction="mean"))))
```

1. Preparing a new configuration file for CoMatch

Creating:

```
./Classification-SemiCLS/  
configs/comatch_ucm.py
```

Modifications: configuration file



```
model = dict(  
    type="resnet18", #config for resnet purposes  
    width=1,  
    in_channel=3,  
    num_class=num_classes,  
    proj=True,  
    low_dim=64, # projection head  
)
```

Using ResNet18 to reduce GPU usage

Proj=True to get the projection head (embeddings)

Other modifications: reduce image size to 60x60 and augment batch size to 32

Not all **models** use same amount of GPU

RN-18

3.5 GB

RN-50

7 GB

WRN-

28-2

12 GB

- Image size: 60x60
- Batch size: 32
- CoMatch
- Labeled to unlabeled ratio: 1-7




04

Model evaluation and inferencing

Analyze the performance of the model via
classification accuracy metrics and embedding
visualization

17:00 - 17:30 (30 min)

Classification accuracy (from logs)



Best model was obtained on the epoch: 200
The best top-1 accuracy corresponds to: 94.47

Case 1:

- Fixmatch on UCM
- 4 labels per class
- WideResNet (Batch=8)

Your directory structure should now resemble the following:

Classification-SemiCLS

results

fm_ucm_wres_x2_b1x8_l840

fm_ucm_wres_x2_b1x8_l840.log

fm_ucm_wres_x2_b1x8_l840.tar

fm_ucm_wres_x2_b1x8_l840.py

- The `.py` file represents the configuration file.
- The `.log` file contains the training output.
- The `.tar` file contains the trained model.

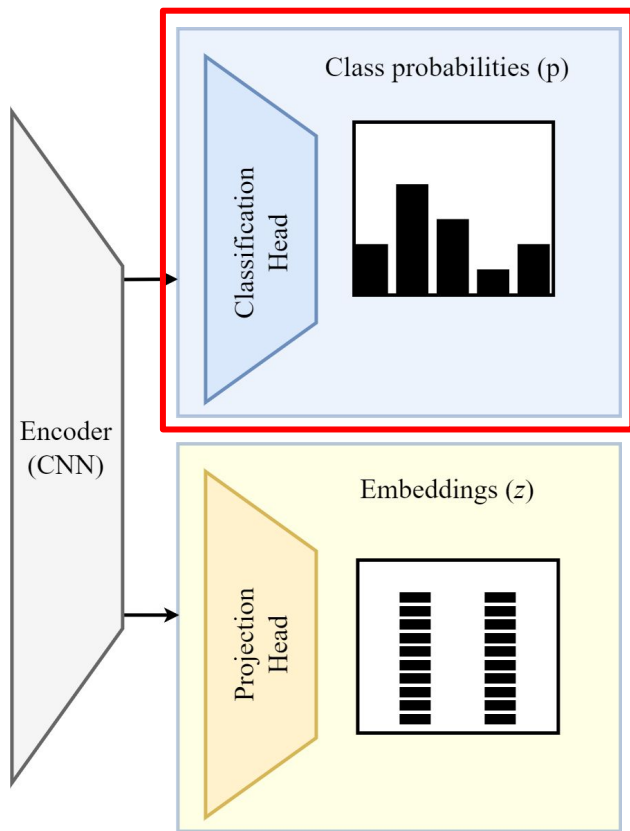
Computational Cost (from logs)

Trainer	Model	Total Parameters	Num Classes Labeled	Mu	Batch size	Training Time
FixMatch	Wideresnet	1.47M	21 4xclass	7	8	2 days, 19:39:18

Data from:

```
./Classification-SemiCLS/results/computational_costs.csv
```

Inference



```
# wideresnet.py (models)
L. 165 for encoder and 1.190 for classification head
```

```
#encoder CNN
feat = self.conv1(x)
feat = self.block1(feat)
feat = self.block2(feat)
feat = self.block3(feat)
feat = self.relu(self.bn1(feat))
feat = F.adaptive_avg_pool2d(feat, 1)
feat = feat.view(-1, self.channels)
```

Conv: Convolution layers

Block: Network blocks with multiple convolutional layers for feature extraction

Relu: Activation functions

BN: Batch normalization layers

```
# output
out = self.fc(feat)
return out
```

fc: fully connected layer followed by a softmax for classification - class predictions.


```
outputs = model(inputs)
```

Classification accuracy (inference)

FixMatch - 5 min inference

Kappa: 0.896



True: golfcourse
Predicted: river



True: harbor
Predicted: intersection



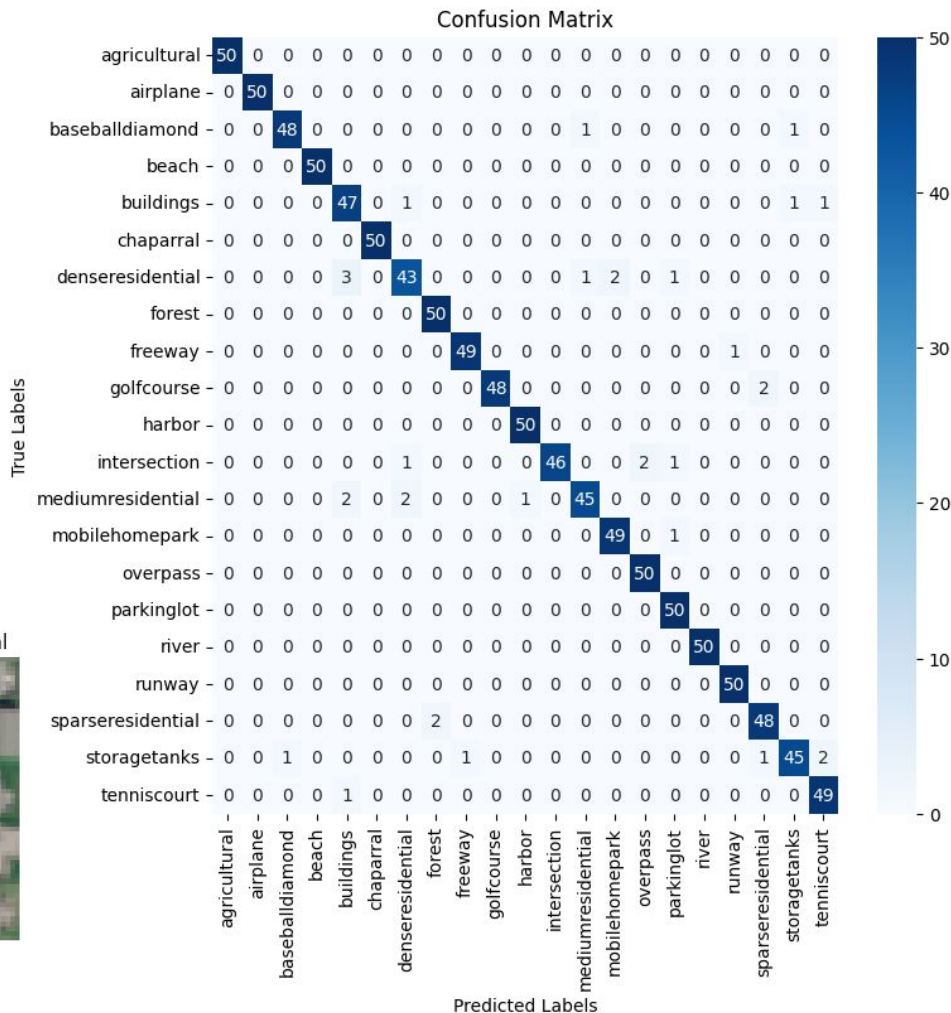
True: forest
Predicted: agricultural



Classification accuracy (inference)

CoMatch - 10 min inference

Kappa: 0.967



True: intersection
Predicted: denseresidential



True: mediumresidential
Predicted: denseresidential



True: mediumresidential
Predicted: denseresidential




Embeddings

```
#--cfg
model = dict(
    proj=True,
```

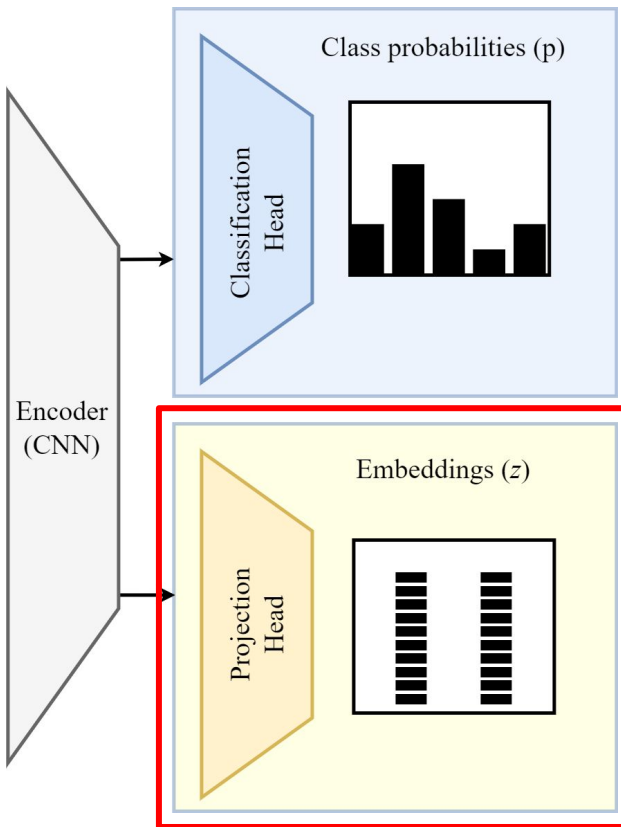
Low-dimensional representation of high-dimensional data.

CNN generate embeddings that capture the underlying structure and relationship of the data.

Projection head: generating feature embeddings for contrastive learning.



Embeddings



L2norm: Normalization layer


Fc1: Fully connected with LeakyReLU activation

Fc2: Another fully connected to reduce dimensionality of the embeddings.

```
#WideResNet.py
# projection head
if self.proj:
    self.l2norm = Normalize(2)
    self.fc1 = nn.Linear(64 * self.widen_factor, 64 *
self.widen_factor)
    self.relu_mlp = nn.LeakyReLU(inplace=True, negative_slope=0.1)
    self.fc2 = nn.Linear(64 * self.widen_factor, self.low_dim)
```


Embeddings

2D or 3D?



Highlight session



01

Introduction to DSSL

16:00 - 16:15 (15 min)

Feedback:

ML and DL

DL RS

CNN

SSL Taxonomy

FixMatch

CoMatch

Highlight session



02

Data & Code

16:15 - 16:30 (15 min)

Feedback:

- Clone GitHub
- Download data
- Sampling strategy
- Train/Test.txt
- Augmentation
- Pipeline

Highlight session



03

Training model

16:30 - 17:00 (30 min)

Feedback:

- Check GPU
- Environment setup
- Code modifications
- Adapt configuration files
- Launch training

Highlight session



04

Evaluate model

17:00 - 17:30 (30 min)

Feedback:

- Extract data from logs
- Classification Accuracy
- Computational Costs
- Inference with Classification Head (class probabilities)
- Inference with Projection Head (embeddings)

Thanks!

Do you have any questions?

isequeir@uji.es



@itzahs



References

Bibliography:

1. Bhavsar, K. (2018, September 14). Understanding your Convolution Network with Visualizations. Towards Data Science. <https://towardsdatascience.com/understanding-your-convolution-network-with-visualizations-a4883441533b>
2. G. Cheng, X. Xie, J. Han, L. Guo and G. -S. Xia, "Remote Sensing Image Scene Classification Meets Deep Learning: Challenges, Methods, Benchmarks, and Opportunities," in IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing, vol. 13, pp. 3735-3756, 2020, [doi: 10.1109/JSTARS.2020.3005403](https://doi.org/10.1109/JSTARS.2020.3005403).
3. X. Yang, Z. Song, I. King and Z. Xu, "A Survey on Deep Semi-Supervised Learning," in IEEE Transactions on Knowledge and Data Engineering, 2022, [doi: 10.1109/TKDE.2022.3220219](https://doi.org/10.1109/TKDE.2022.3220219).
4. Sharma, A. (2019, June 12). Complete Guide to Data Augmentation for Computer Vision. Towards Data Science. <https://towardsdatascience.com/complete-guide-to-data-augmentation-for-computer-vision-1abe4063ad07>.
5. Shorten, C., Khoshgoftaar, T.M. A survey on Image Data Augmentation for Deep Learning. *J Big Data* 6, 60 (2019). <https://doi.org/10.1186/s40537-019-0197-0>
6. Kihyuk Sohn, David Berthelot, Nicholas Carlini, Zizhao Zhang, Han Zhang, Colin A Raffel, Ekin Dogus Cubuk, Alexey Kurakin, and Chun-Liang Li. Fixmatch: Simplifying semi-supervised learning with consistency and confidence. *NeurIPS*, 33, 2020.
7. Yi Yang and Shawn Newsam, "Bag-Of-Visual-Words and Spatial Extensions for Land-Use Classification," ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems (ACM GIS), 2010.
8. F. Yang et al., "Class-Aware Contrastive Semi-Supervised Learning," 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), New Orleans, LA, USA, 2022, pp. 14401-14410, [doi: 10.1109/CVPR52688.2022.01402](https://doi.org/10.1109/CVPR52688.2022.01402).
9. S. Zagoruyko and N. Komodakis, "Wide residual networks," in Proceedings of the British Machine Vision Conference (BMVC), 2017, pp. 87.1-87.12, [doi: 10.22024/UniKent/01.02.69550](https://doi.org/10.22024/UniKent/01.02.69550), arXiv:1605.07146.

Presentation credits:

1. Template by **Slidesgo**
2. Icons by **Flaticon**
3. Infographics and images by **Freepik**

